



AMERICAN
UNIVERSITY
OF BEIRUT

Fall 2025

EECE 503P/798S: Agentic Systems

C3 - Introduction to Agents



Lesson Objectives

- Define what an **AI agent** is and identify real-world applications
- Explain the **Reinforcement Learning (RL)** setup of an agent
- Understand what makes a system **agentic**
- Break down the **core components** of an LLM-based agentic system (tools, memory, planning, etc.)
- Recognize **key challenges** in deploying agentic systems (e.g., hallucinations, cost, debugging)



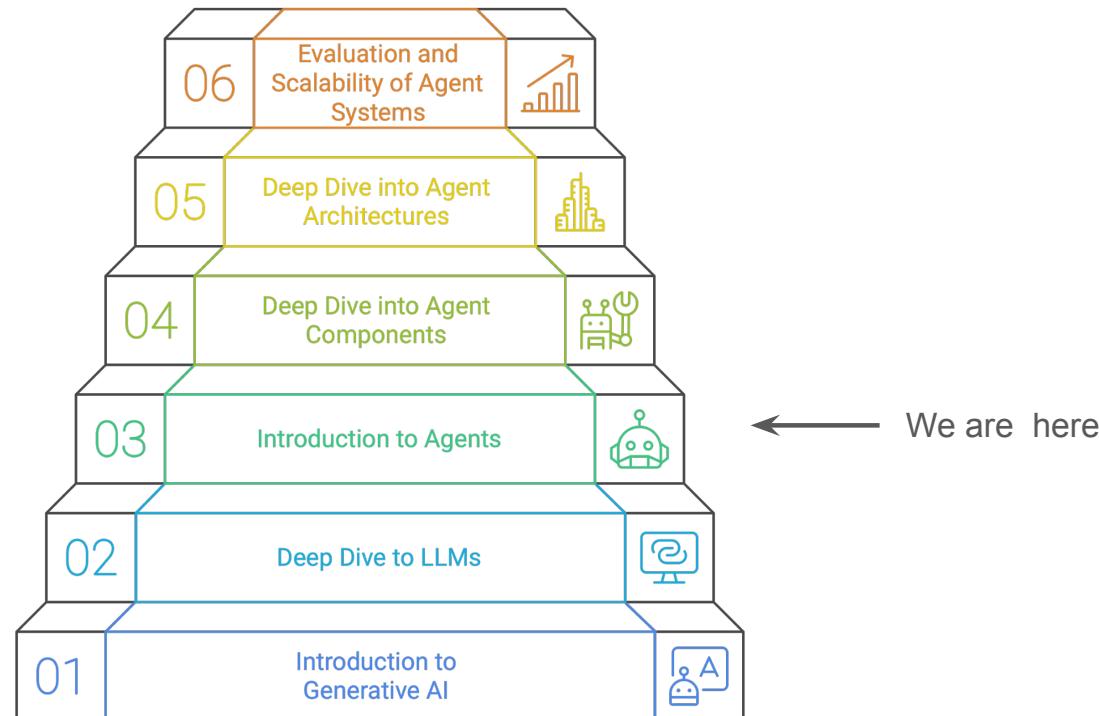
Introduction

This chapter is inspired by the Udemy course: [The Complete Agentic AI Engineering Course \(2025\)](#) by Ed Donner.

You'll find everything you need to setup your environment in this Github repo: [agents by Ed Donner](#).



Course Timeline



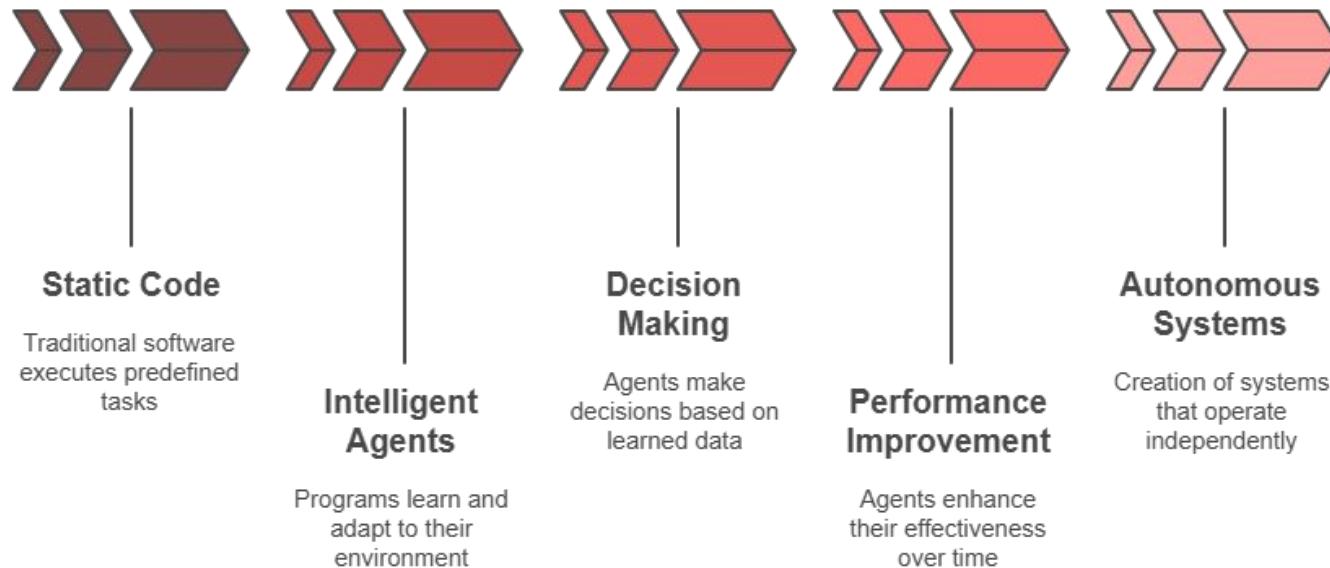


AMERICAN
UNIVERSITY
OF BEIRUT

Introduction



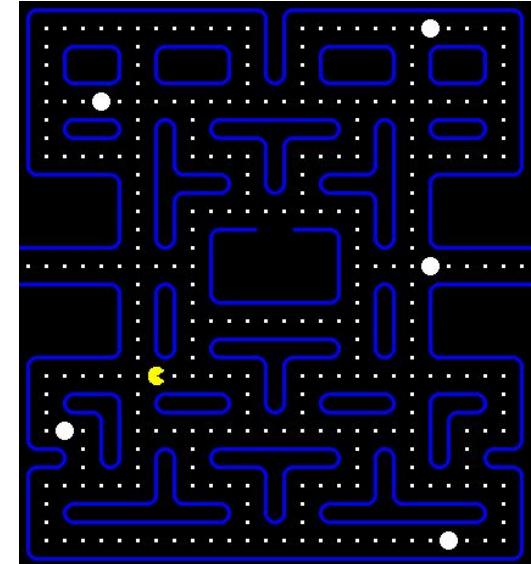
From Static Code to Intelligent Agents





Agents are Not New

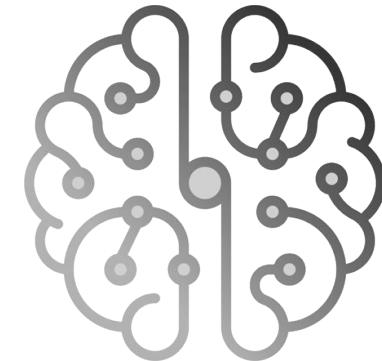
- The concept of AI agents **predates** LLMs by decades
- **Early agents (1990s–2000s):**
 - Based on **symbolic AI, rule-based reasoning, or classical planning**
 - Used in **robotics, game agents, and expert systems**
- Focused on **perception + action** within well-defined **environments**





Why Now?

- Modern LLMs unlocked true agentic behavior
- What LLMs Bring:
 - Language **understanding** and **generation**
 - **Chain-of-thought** reasoning
 - **Generalization** across domains
 - Multi-step **planning** with **feedback**





The Agent Need



Business Needs

Businesses require intelligent automation, not mere scripts.



Developer Needs

Developers desire systems capable of learning and adapting.



User Needs

Users need agents understanding goals, not just command execution.



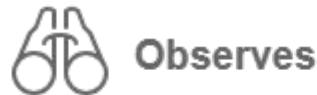
AMERICAN
UNIVERSITY
OF BEIRUT

Definition and Applications of Agents



What is an Agent?

- An agent is an **autonomous system** that:



Observes



Reasons

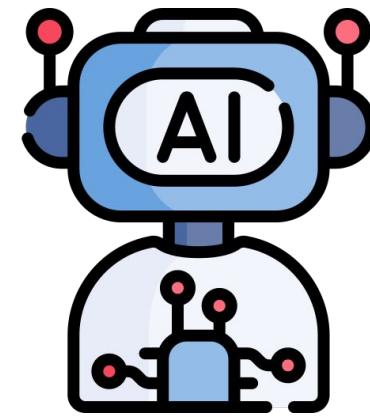


Acts toward achieving goals



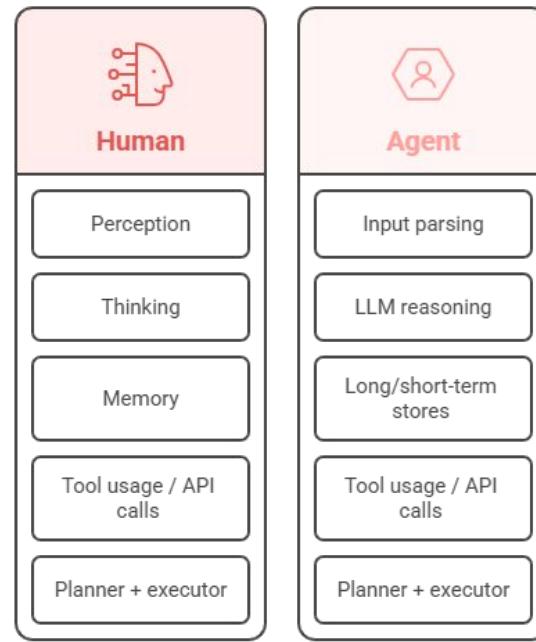
What is an Agent?

- **Agent ≠ Chatbot:**
 - Not just conversation, but action
- **Agent ≠ Script:**
 - Not just automation, but goal-driven reasoning
- **Varying levels of autonomy**
 - from tool caller to planner





How Agents Resemble Human Thinking





Is it an Agent?

- Is a smart thermostat an agent?
- Smart Light Timer?
- Scheduled Email Sender?
- A Roomba (robot Vacuum Cleaner)?



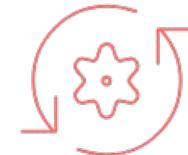
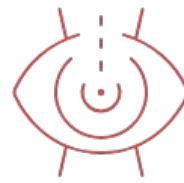
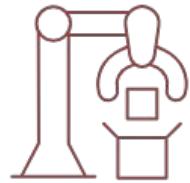


Is it an Agent?

System	Autonomy	Perception	Action	Goal-Oriented	Agent?
Smart Thermostat	✓	✓	✓	✓	✓
Roomba	✓	✓	✓	✓	✓
Smart Light Timer	✓	✗	✓	✗	✗
Scheduled Email Sender	✓	✗	✓	✗	✗



Key Characteristics of Agents



Autonomy

Operates without human intervention

Perception

Observes environment/state

Action/Decision-making

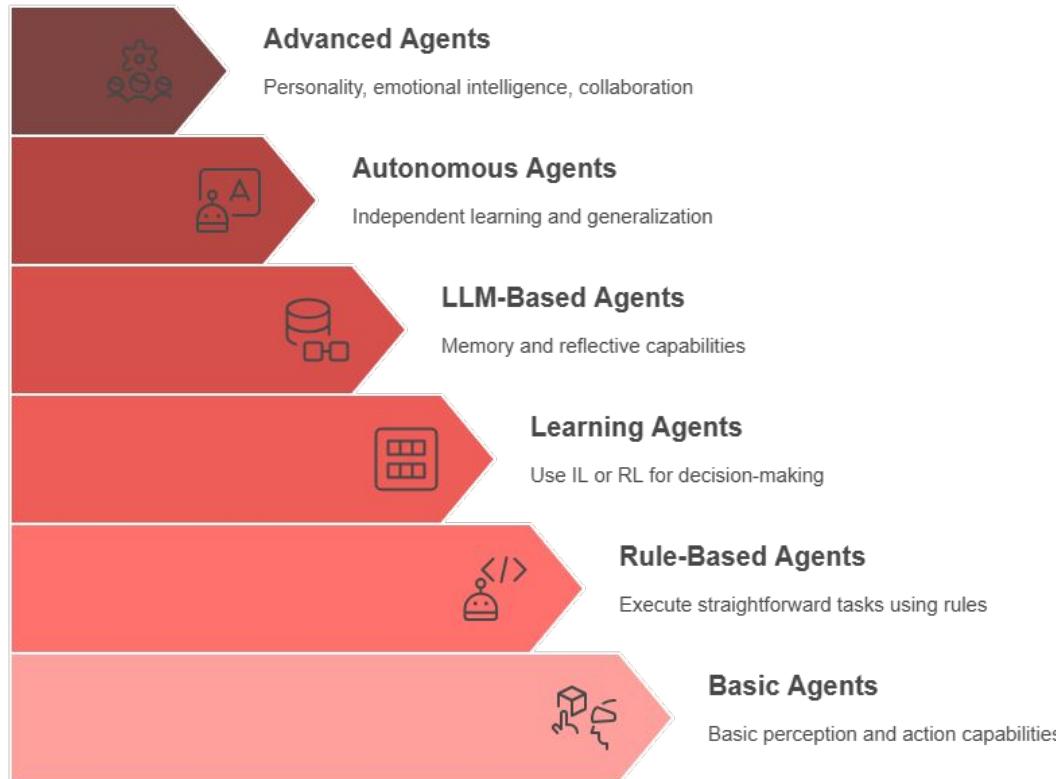
Takes steps to change state

Goal-Oriented

Acts to achieve objectives



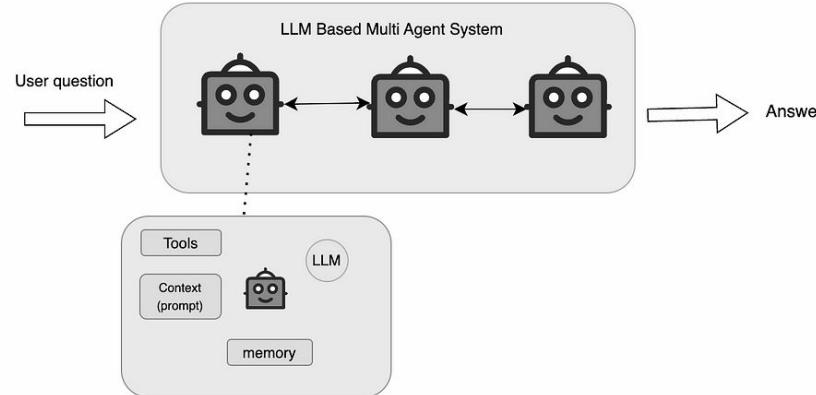
Categorization Based on Task Complexity





AI/LLM-Based Agents

- Combine the **classic** agent loop with **natural language understanding**
- Use **LLMs** (e.g., ChatGPT, Claude, Gemini) as **reasoning engine**
- Can **interpret goals, choose tools, and plan actions** dynamically
- **Examples:** LangChain agents, AutoGPT, Devin, CrewAI

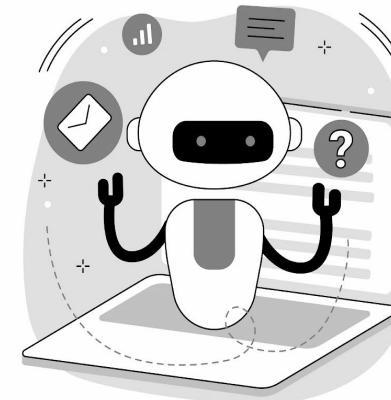




AI Agents Definition

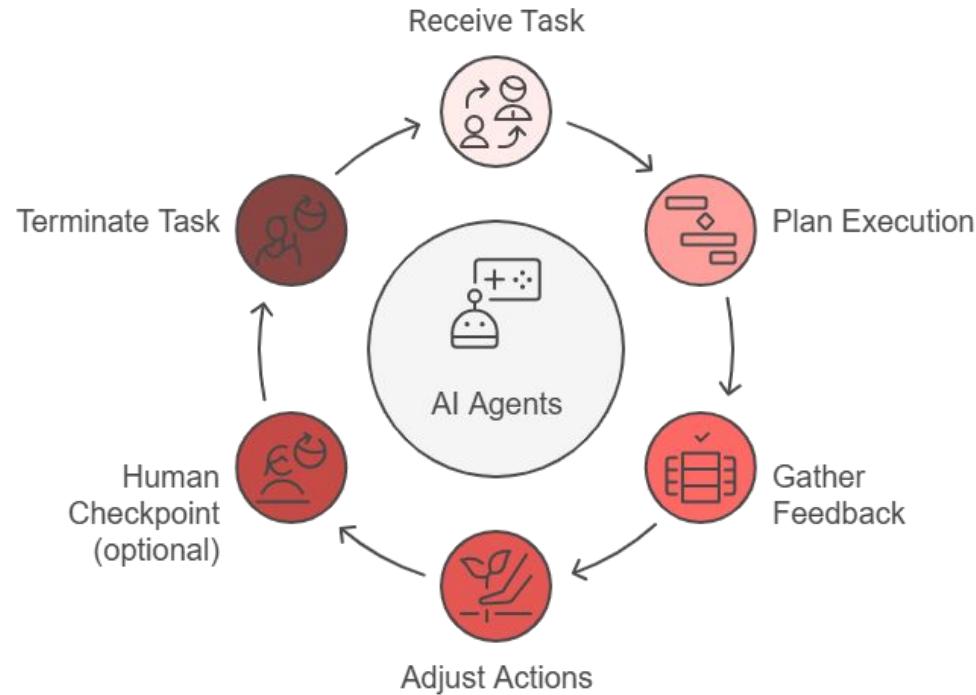
AI agents are programs where LLM **outputs** govern the **workflow**

1. **Multiple Calls to LLMs**
2. **LLMs with the Ability to Use Tools**
3. **An Environment Where LLMs Interact**
4. **A Planner Coordinating Activities**
5. **Autonomy**





Core Lifecycle of an Agent



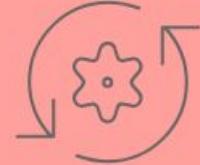


When to use Agents



Task Complexity

Agents are useful for autonomous tasks with complex decisions.



Repetitive Workflows

Agents automate frequent and routine workflows efficiently and accurately.



Unstructured Data

Agents process documents, conversations, and non-tabular data effectively.



Customer Support Agents

- Blend **conversational flow** with **action-taking** via integrated tools
- Access **user data, order history, and knowledge base articles**
- **Automate** actions (e.g., refunds, ticket updates)
- Human oversight ensures **trust and quality**





Coding Agents

- Solve programming problems **autonomously**
- Use **automated tests** for verification and feedback
- Operate in structured, well-defined **environments** (e.g., GitHub issues)
- Human review ensures **alignment** with broader system goals





AMERICAN
UNIVERSITY
OF BEIRUT

The RL setup of agents



Find the Treasure

No instructions. No map.

Just one goal: “**Find the treasure.**”

- You try going left — *dead end*
- You try going right — *trap!*
- You learn from failure
- You try again — *you get better*
- Eventually, you succeed
- That’s not luck — ***that’s learning***





What Is Reinforcement Learning?

RL is like teaching a pet tricks — give it treats (**rewards**) when it does the right thing!

Agents **learn** what to do by:

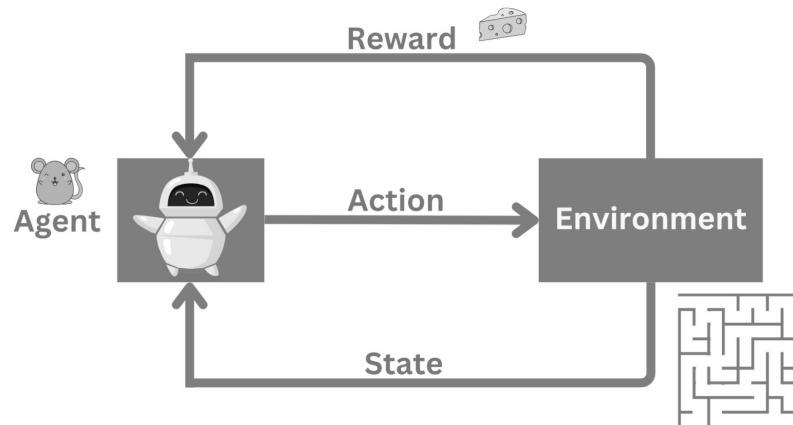
- Trying actions
- Seeing what works (reward)
- Trying again (better next time!)





What Is Reinforcement Learning?

- An agent is a **decision-making** entity
- It interacts with an **environment** through **trial and error**
- Its goal is to **maximize rewards** over time
- Inspired by how humans and animals **learn** from experience





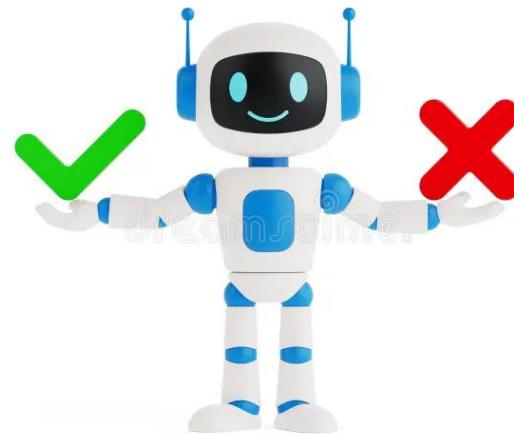
The Agent-Environment Loop





Core Components of RL- Agent

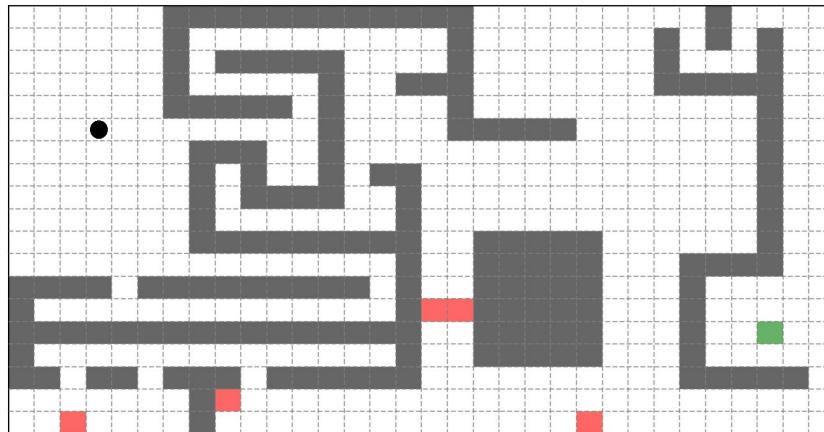
- The AI system that makes **decisions** in order to achieve a **goal**
- Can be a robot, game character, chatbot, or software





Core Components of RL- Environment

- Everything the agent **interacts with** and receives **feedback** from
- Provides the current **state** and **reward**
- Can be **physical** (robot's world) or **digital** (video game)





Core Components of RL- State

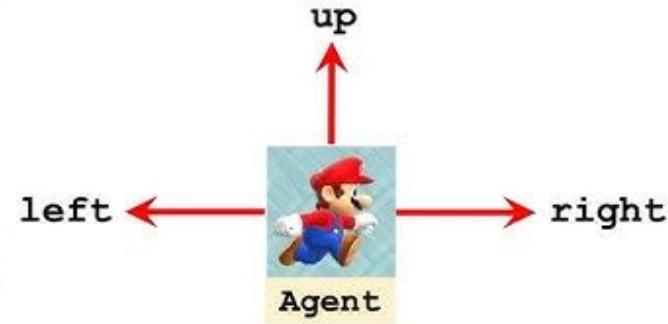
- The agent's view of the **current situation** in the environment
- Can be **full** (chess board) or **partial** (robot with limited sensors)
- Influences which **action** is chosen
- Often represented as **vectors, images, or structured data**





Core Components of RL- Action

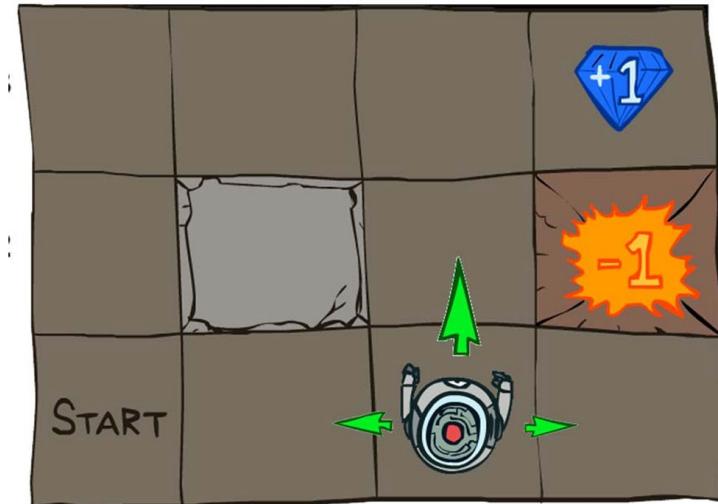
- A **decision** or **behavior** the agent performs at a **specific time**
- Chosen based on the **current state** and **policy**
- **Examples:** Move left, turn off light, buy stock, answer question
- Some actions may be **better** than others depending on the context





Core Components of RL- Reward

- A signal telling the agent how **good** or **bad** its action was
- **Example:**
 - Positive = good action (e.g., +10 for success)
 - Negative = bad action (e.g., -10 for failure)
- The goal is to **maximize** the **cumulative reward**





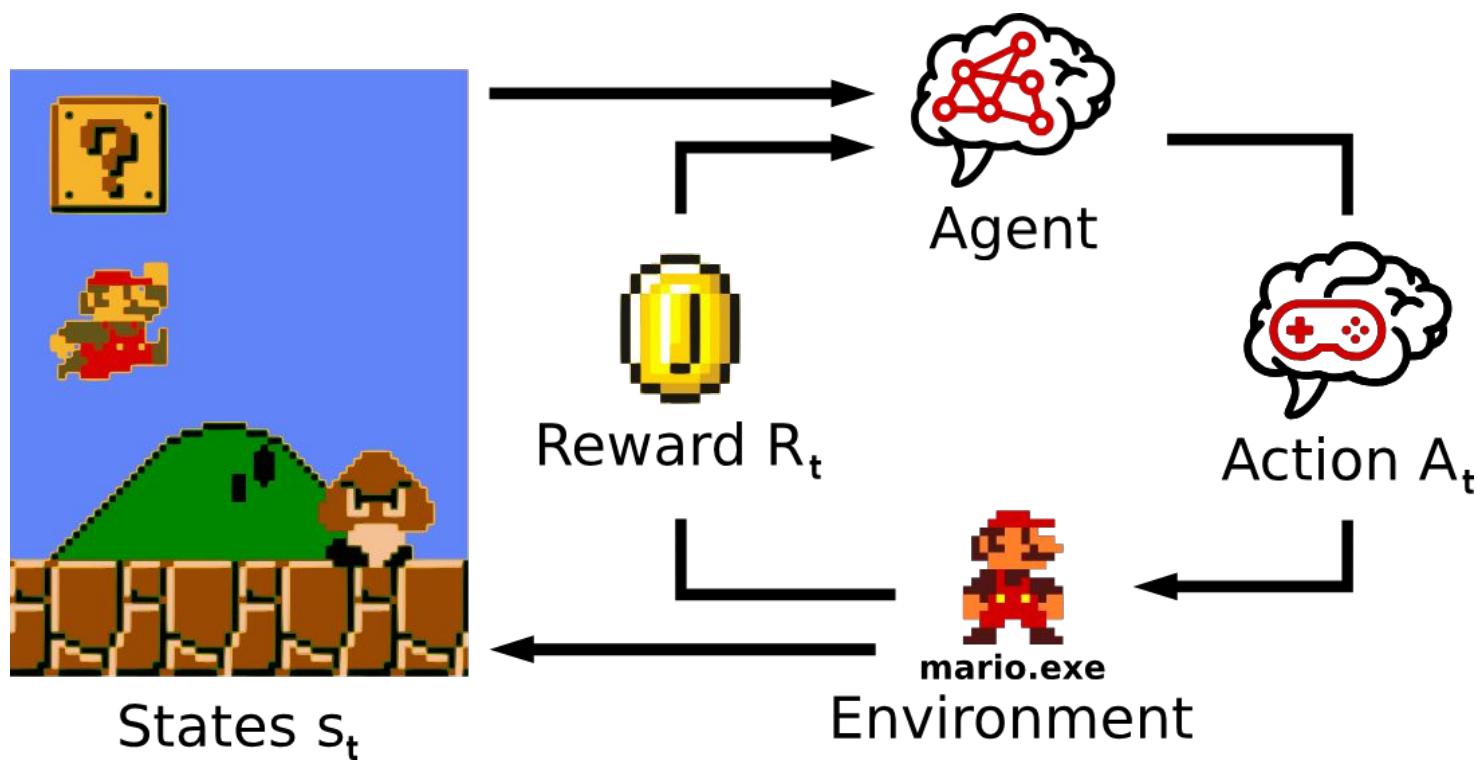
Core Components of RL- Policy

- Mapping from **state to action**: "In this situation, do this."
- The agent's learned or programmed **decision plan**
- **Updated over time** as the agent learns
- In RL, often represented as $\pi(a|s)$ = **probability of action a in state s**





Reminder: Putting it together





Example: Smart Vacuum Cleaner



State:
room is dirty



Action:
clean or move



Reward:
+ 1 for cleaning



Policy:
If room is dirty → Clean
If room is clean → Move



RL Vs. LLM Agents

Feature	RL Agent	LLM Agent
 Learning Method	Rewards & Trials	Prompts, Tools, Memory
 Policy	Trained Function	Emergent from LLM Behavior
 Tool Use	Rare	Frequent
 Flexibility	Learns from Data	Generalizes via Pretraining



Challenges in RL

Pros

- Handles complex tasks
- Learns from data



Cons

- Slow learning
- Delayed feedback
- Repetitive behavior
- Debugging challenges





AMERICAN
UNIVERSITY
OF BEIRUT

Hands_on: Cleaning Agent Navigation in a Grid World



AMERICAN
UNIVERSITY
OF BEIRUT

Agentic Systems

What does “agentic” even mean?

- Derived from "agent" — an entity capable of acting **independently**
- Refers to systems that exhibit **autonomy, goal pursuit, and decision-making**
- **Example:**
 - **Non-Agentic AI:** A calculator (responds to direct input)
 - **Agentic AI:** A personal assistant (understands your intent and takes steps to fulfill it)





The Rise of Agentic Systems

1 Agents effectively utilize tools for task completion.

Tool Usage

2 Agents access and retrieve relevant information efficiently.

Memory Access

3 Agents strategically plan and execute tasks effectively.

Task Planning

4 Agents collaborate seamlessly with other agents.

Agent Collaboration

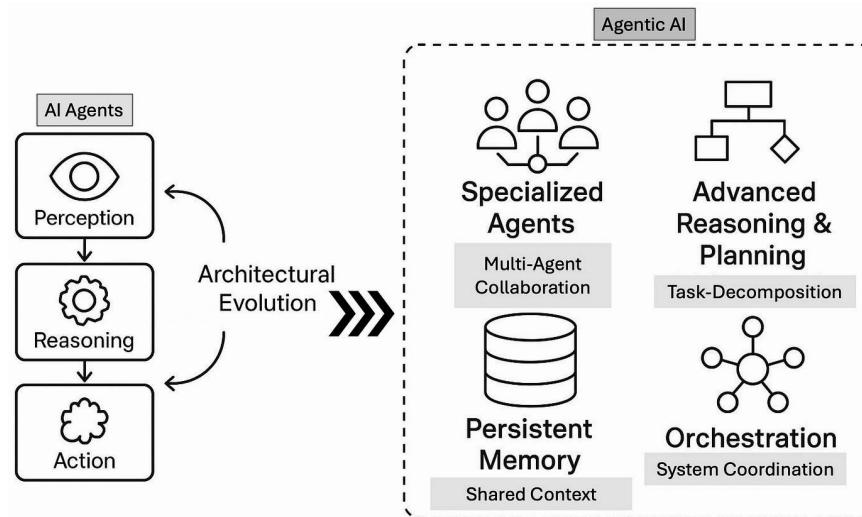
Agentic Systems





Agentic Systems

- Systems that coordinate multiple **autonomous** agents
- May include shared **memory**, **planning** components, or **role** assignment
- Can **divide** and **conquer** tasks (task planning + execution + verification)
- Often involve **feedback loops** and **multi-step tool** use

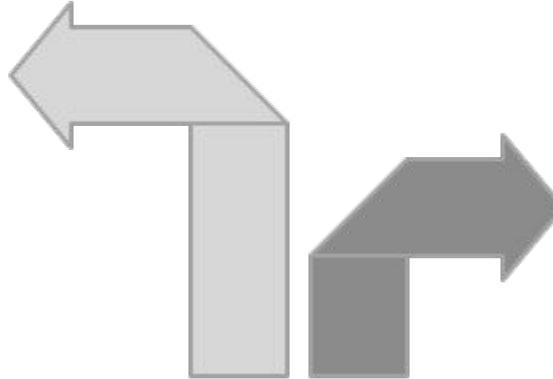




Example: Fraud Detection

Agentic Workflow

The AI detects an unusual transaction, sends a notification to the customer, and verifies the purchase through a quick app interaction

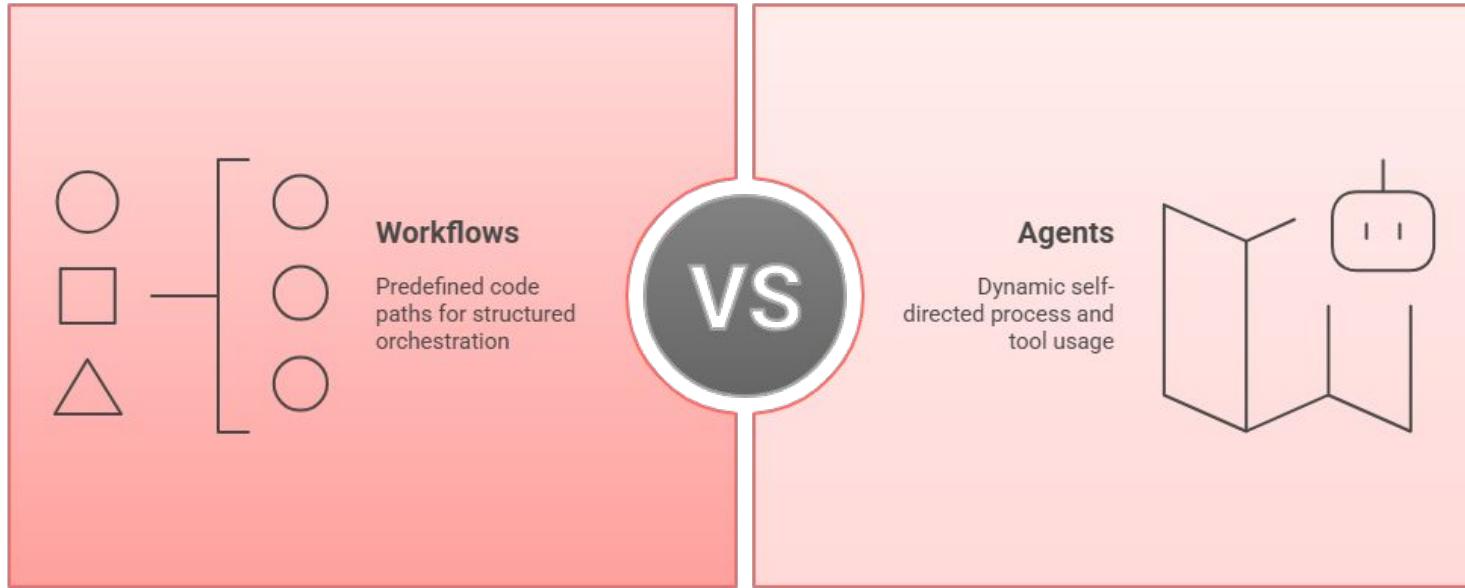


Non-Agentic Workflow

A customer's credit card purchase is flagged as fraudulent, requiring them to call the bank and verify the transaction

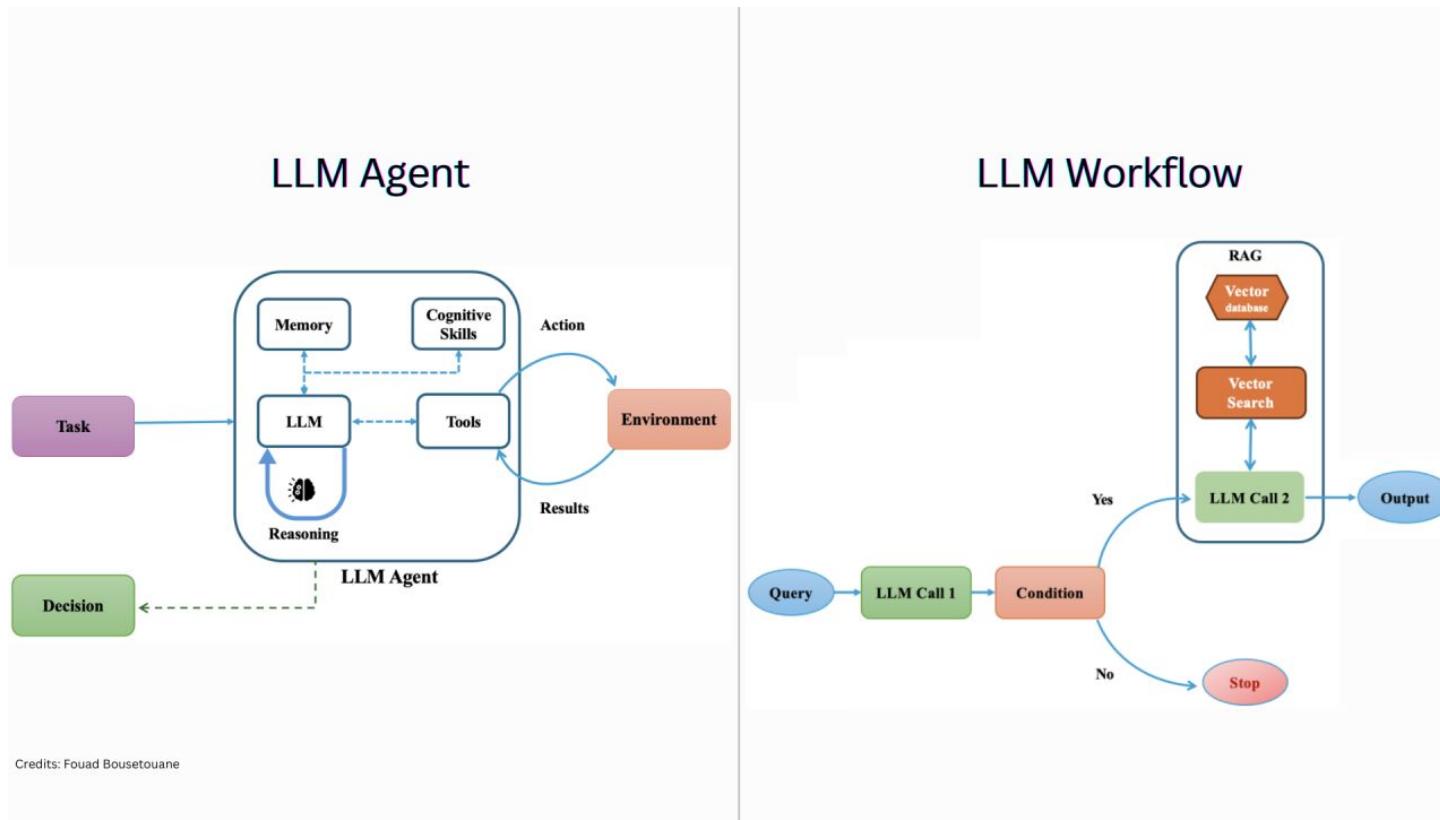


Agentic Systems: Workflows and Agents





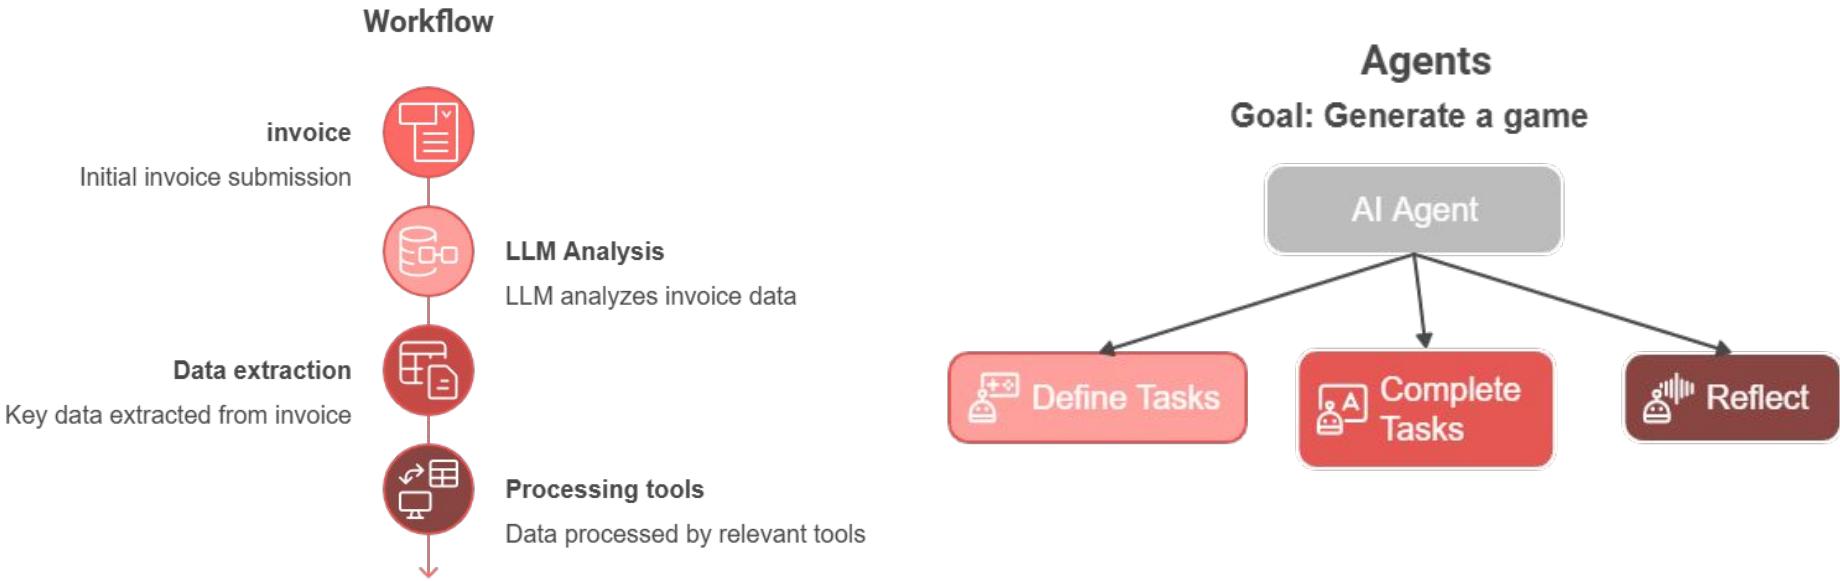
Agentic Systems: Workflows and Agents



Credits: Fouad Bousetouane



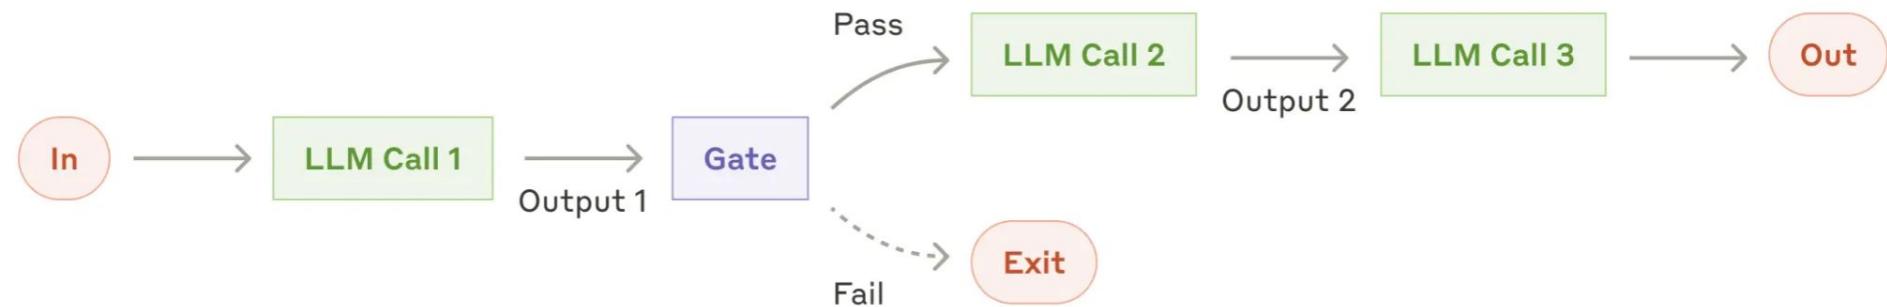
Agentic Systems: Workflows and Agents





Workflow Design Patterns: Prompt Chaining

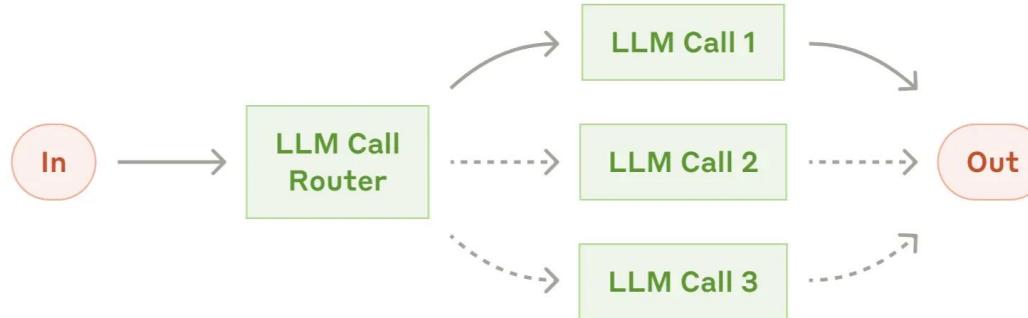
- Breaks a **complex task** into a sequence of **smaller LLM steps**
- Use when tasks can be cleanly **decomposed** into fixed subtasks
- **Example:** Write document outline → Validate outline → Generate document





Workflow Design Patterns: Router

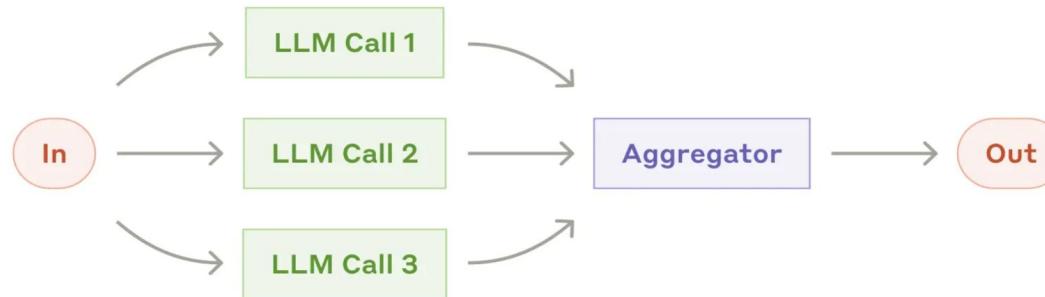
- Direct an input into **specialized subtasks**
- Use to **separate concerns** and **improve performance** per category
- **Example:** route customer queries general vs. refund vs. tech support to different workflows





Workflow Design Patterns: Parallelization

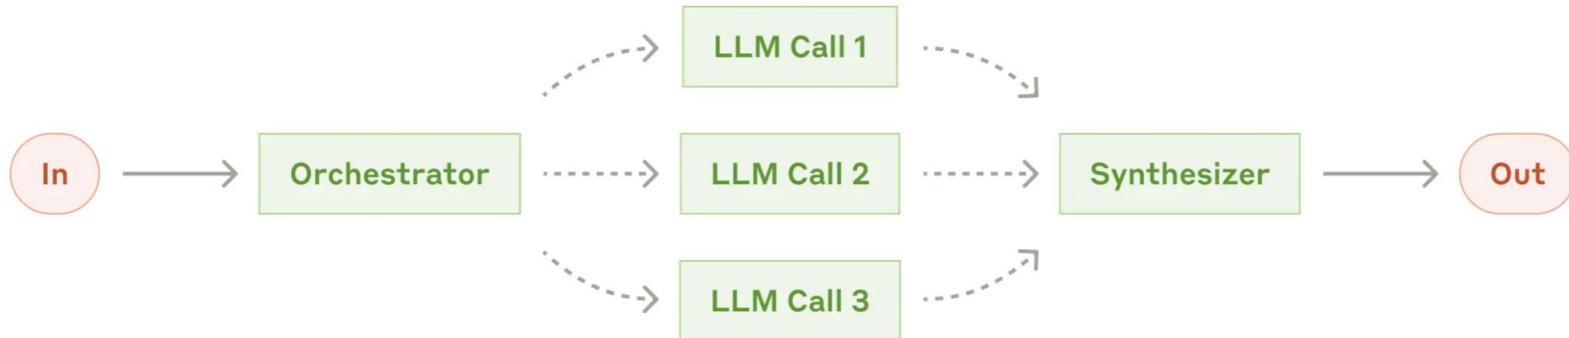
- Breaking down tasks and running multiple subtasks **concurrently**
- **Types:**
 - **Sectioning** – Split task into independent parts, run in parallel (Each LLM scores a different eval metric)
 - **Voting** – Run the same task multiple times, then combine outputs (Multiple reviews of code for bugs)
- Use when tasks can be **parallelized**





Workflow Design Patterns: Orchestrator-workers

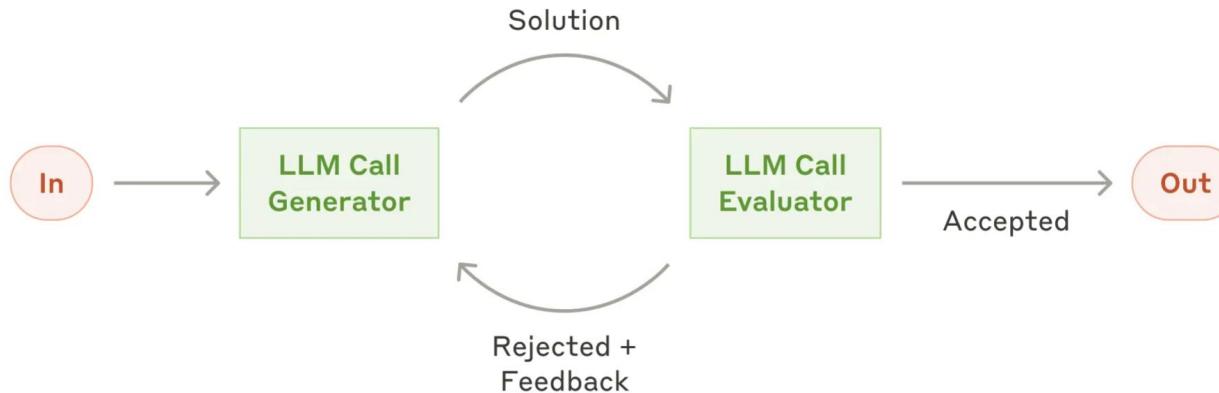
- A central **orchestrator LLM** dynamically breaks down a task, delegates to **worker LLMs**, then **combines** their outputs
- Use when **subtasks vary** depending on the input
- **Example:** Research agents that gather, filter, and synthesize information from diverse sources





Workflow Design Patterns: Evaluator-optimizer

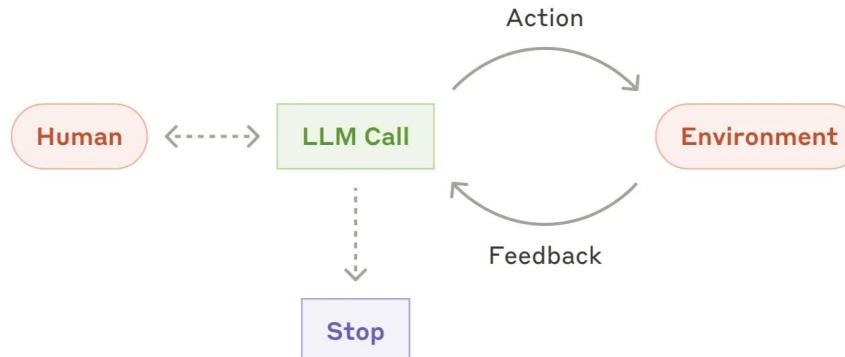
- LLM output is **validated** by another
- Use when **clear evaluation criteria** exist and **iterative improvement** adds real value
- **Example:** Search tasks refined through repeated evaluation and reruns





Agents

- **When to Use:**
 - Open-ended tasks with **unpredictable steps**
 - Tasks that need **reasoning, adaptation, and tool use**
- **Needs guardrails, testing, and clear tool documentation**
- **Example:** Coding agent editing multiple files from a task description





Agentic AI Framework

Langraph

AutoGen

OpenAI Agents
SDK

Crew AI

No framework

MCP



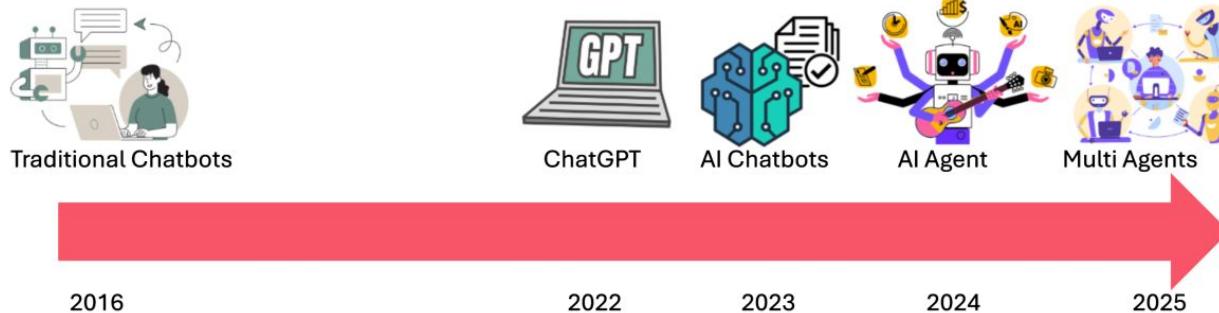
AMERICAN
UNIVERSITY
OF BEIRUT

Core Components of an LLM-based Agentic System



The Evolution of LLM Agents

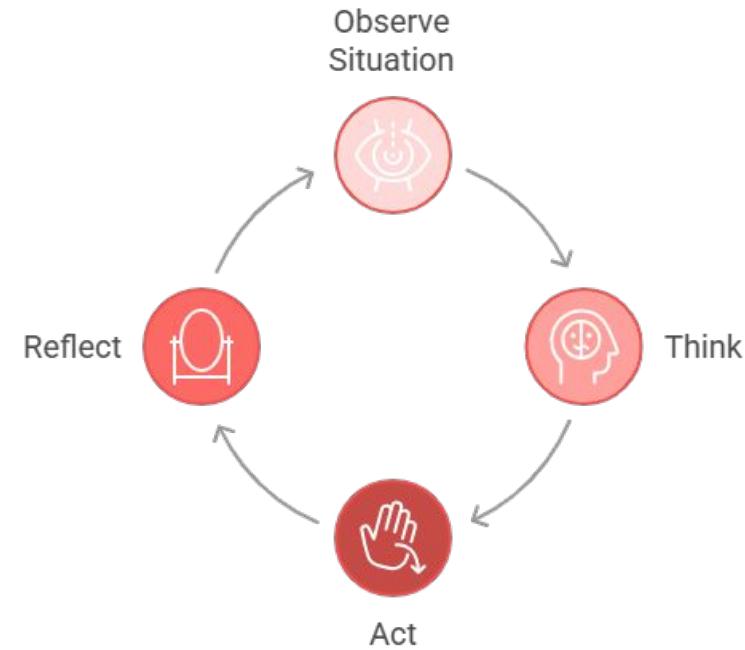
- Initially, LLMs were used in **static settings** (answering questions, generating text)
- Researchers began embedding LLMs into **agent loops**
- Gave rise to the first **LLM-based autonomous agents**
- These agents transitioned from passive responders to **adaptive cognitive engines**





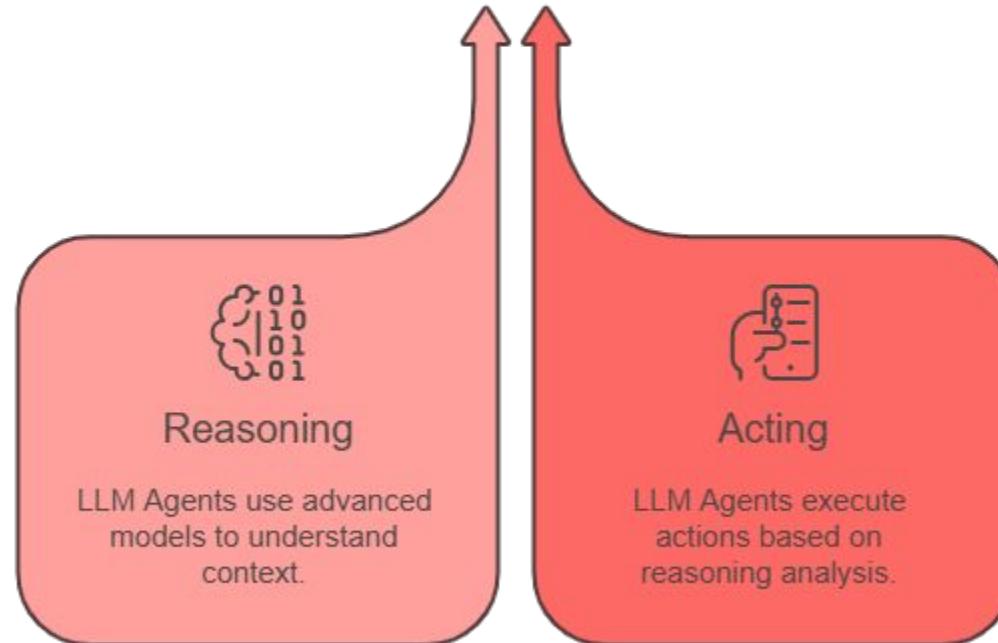
The Agent Loop Revisited

- LLM is not a **one-off responder**, but part of a **feedback loop**
- Each **component** supports one or more **stages** of this loop



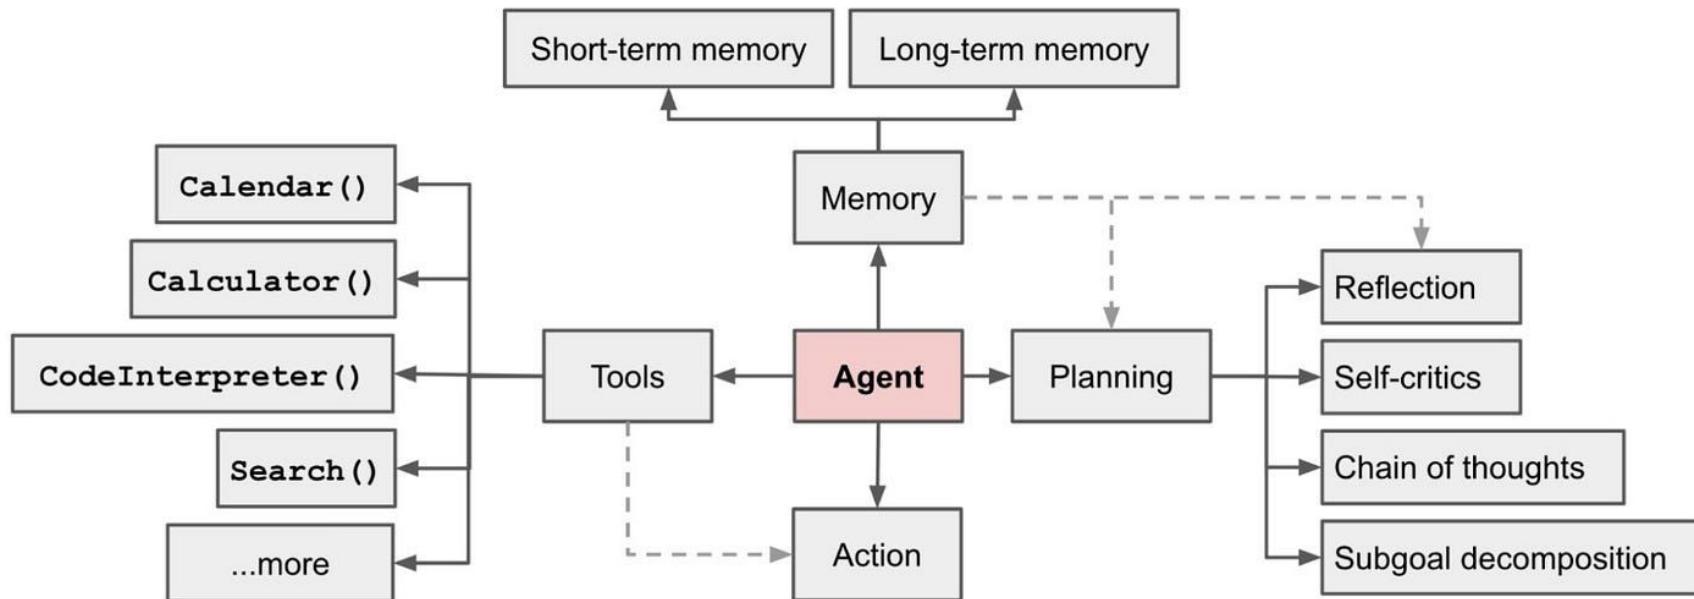


How Do LLM Agents Work?





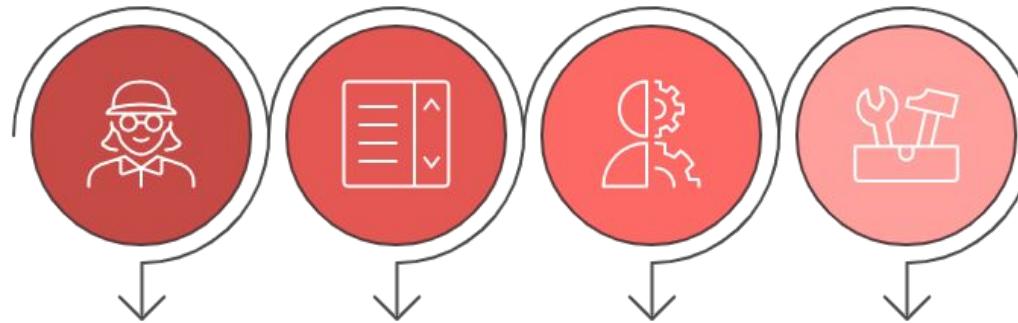
Core Components of LLM-based Agents





Component 1: LLM-based Agent

- When LLMs are given a:



Specific Persona

The LLM is given a specific persona to embody.

Set of Instructions

The LLM receives a set of instructions to follow.

Thinking Mechanisms

The LLM is equipped with thinking mechanisms.

Tools

The LLM is provided with tools to aid its tasks.



Persona — The Agent's Identity

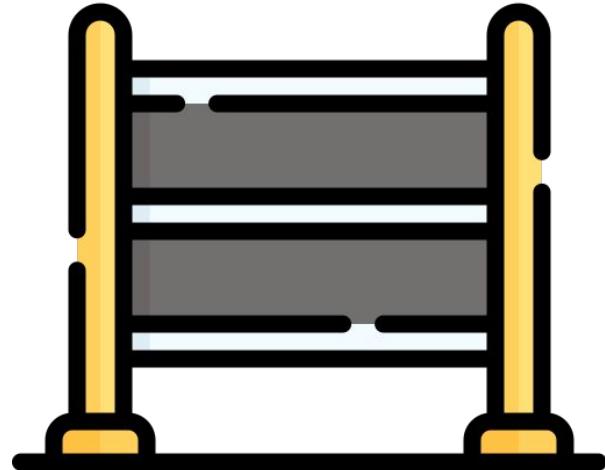
- Defines the **role** the agent plays (e.g., teacher, developer, assistant)
- Shapes **tone, style, and behavior**
- **Examples:**
 - “You are a friendly tutor helping a student.”
 - “You are a cybersecurity expert detecting threats.”





Instructions — Guardrails for Behavior

- Direct the agent's **actions** and **boundaries**
- Include **goals, forbidden actions**, ..
- Act as an **operational manual** for the LLM
- **Examples:**
 - “Avoid making assumptions; ask for clarification.”
 - “Do not access external APIs without permission.”





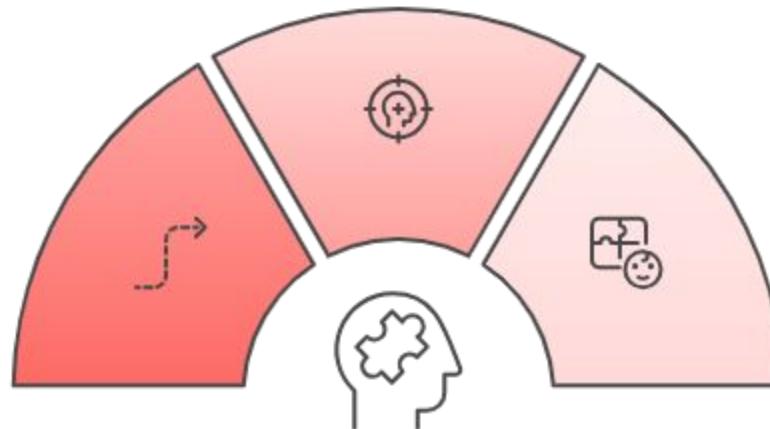
Component 2: Planning/Thinking Mechanisms

Task Decomposition

Simplifies complex tasks by dividing them into manageable subgoals.

Step-by-Step Reasoning

Facilitates a methodical approach to problem-solving.

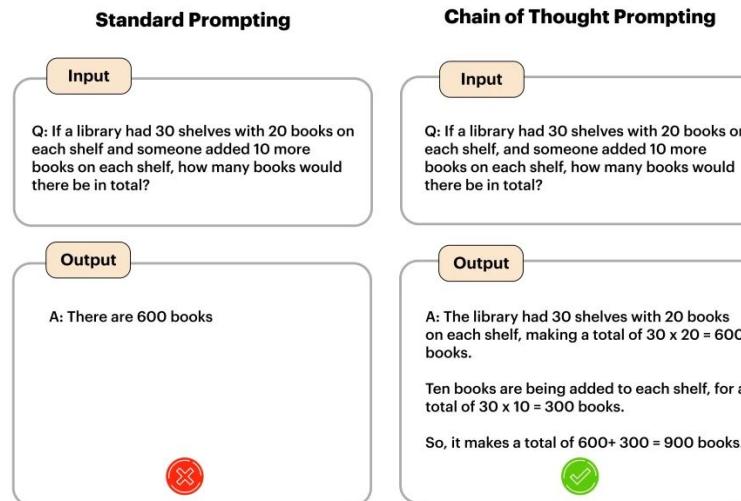


Enhanced Outcomes

Improves the precision, clarity, and effectiveness of actions.

Planning Techniques: Chain of Thought Reasoning

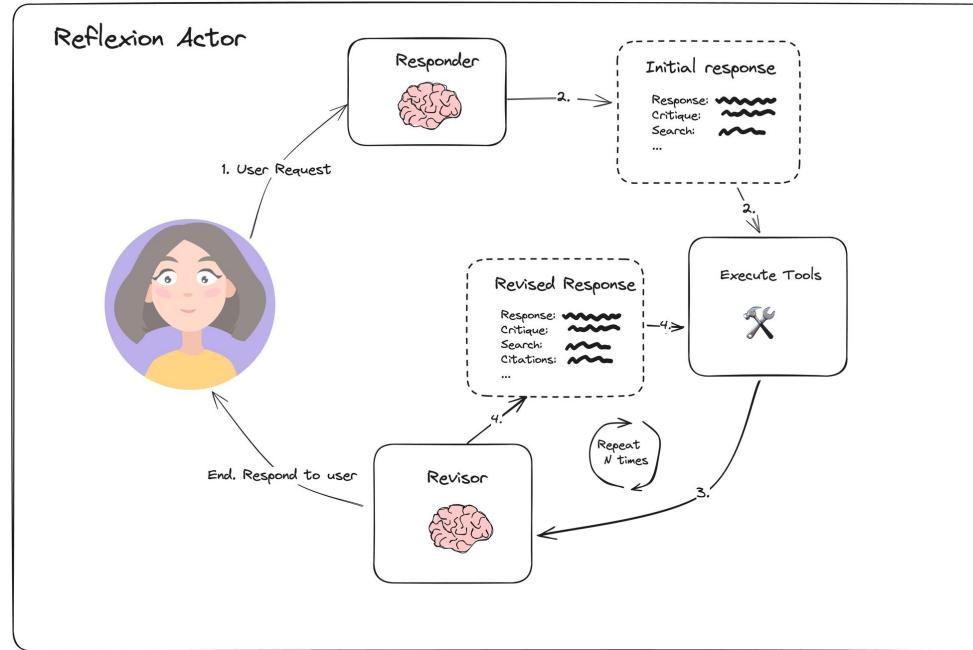
- The agent **explains its reasoning** step-by-step
- Encourages **transparency** and better **problem-solving**





Planning Techniques: Reflection & Self-Criticism

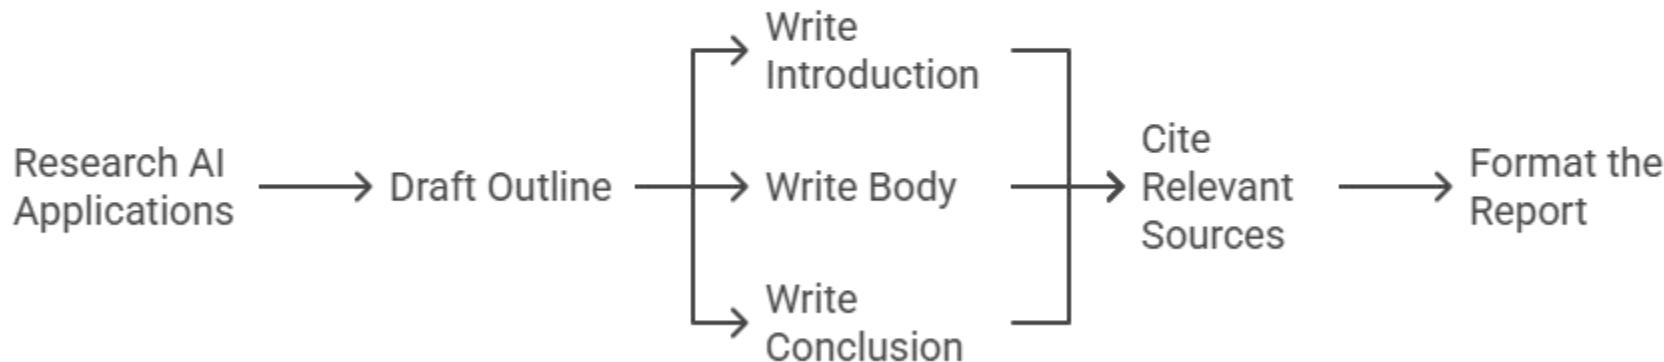
- Agent **evaluates** its own performance
- Detects possible mistakes and **rethinks** answers





Planning Technique: Subgoal Decomposition

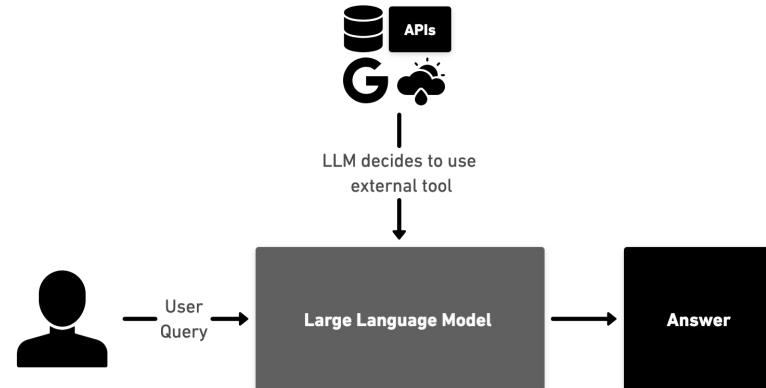
- Agent **breaks** complex goals into **smaller manageable tasks**
- Each subtask can be delegated to **tools** or other **agents**
- Critical for long-term, multi-step **workflows**





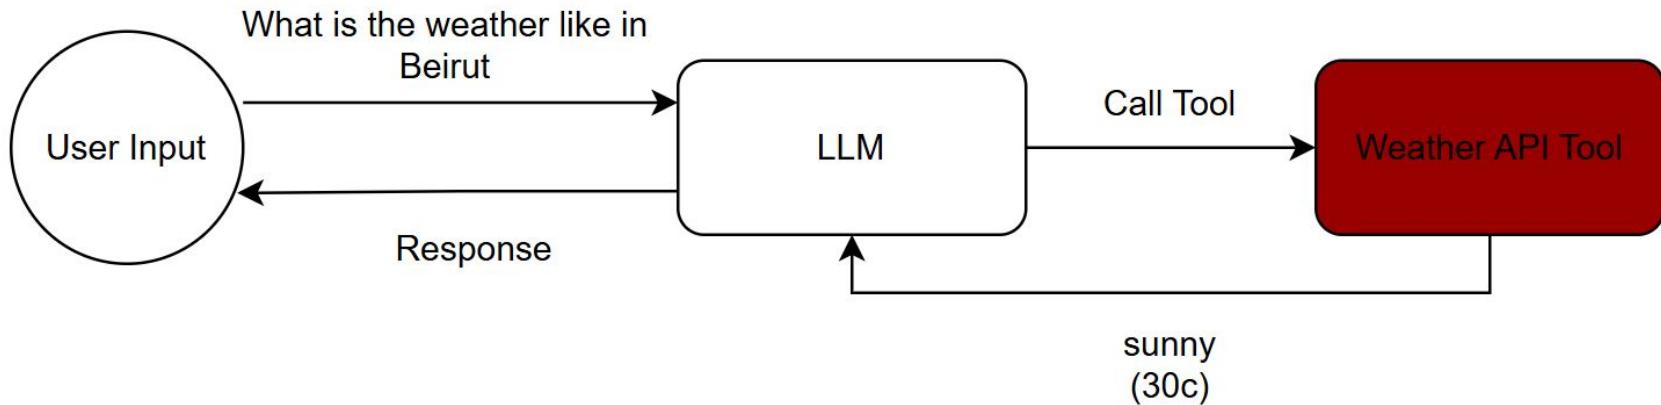
Component 3: Tools

- External functions or APIs the agent can invoke to **solve problems** it cannot handle with text generation alone
- Give LLMs the ability to carry **actions**
- Used for precision tasks (e.g., search, code, math)





Tools Example





Component 4: Action

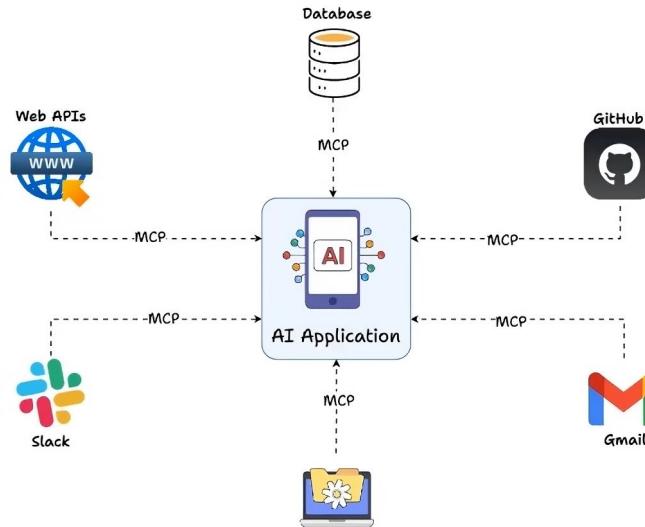
- This is where decisions become **reality** — the agent calls a tool or emits a response
- **Prompts (or fine-tuning)** guide tool use
- **Action schemas** define valid **tools** and **formats**
- **Model Context Protocol (MCP)** standardizes how LLMs call tools





What is Model Context Protocol (MCP)?

- A standard interface for LLMs to **call tools**
- Defines how **models declare, request, and execute actions**
- Supports **tool schemas, parameters, and expected output formats**
- Helps avoid **ambiguous or inconsistent tool calls**





Before vs After Model Context Protocol (MCP)

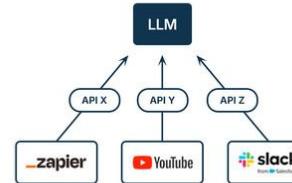
● Before MCP

- No standard way to describe tools
- Each framework had its own tool format
- Difficult to reuse tools across agents or platforms

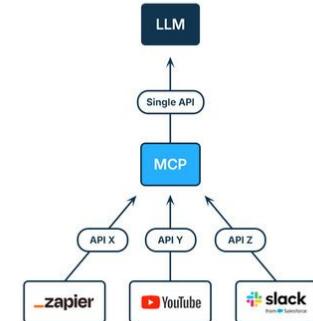
● After MCP

- Tools defined using standard **JSON schemas**
- LLMs know **what inputs** to expect and how to **format outputs**
- Actions can be **validated, parsed**, and **executed automatically**

Before MCP

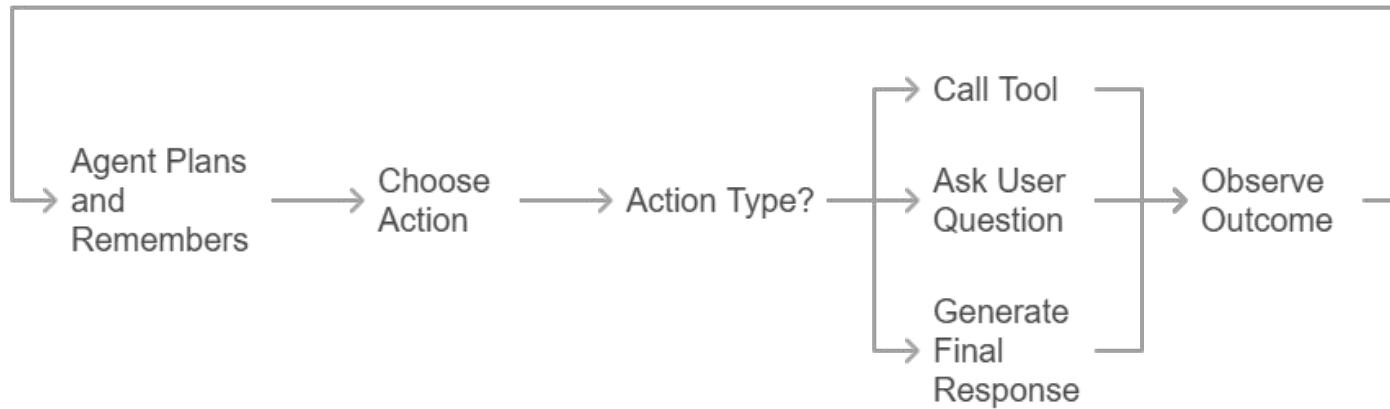


After MCP





Action Example





Hands_on: Experimenting_with_LLM_Configurations _and_Personas



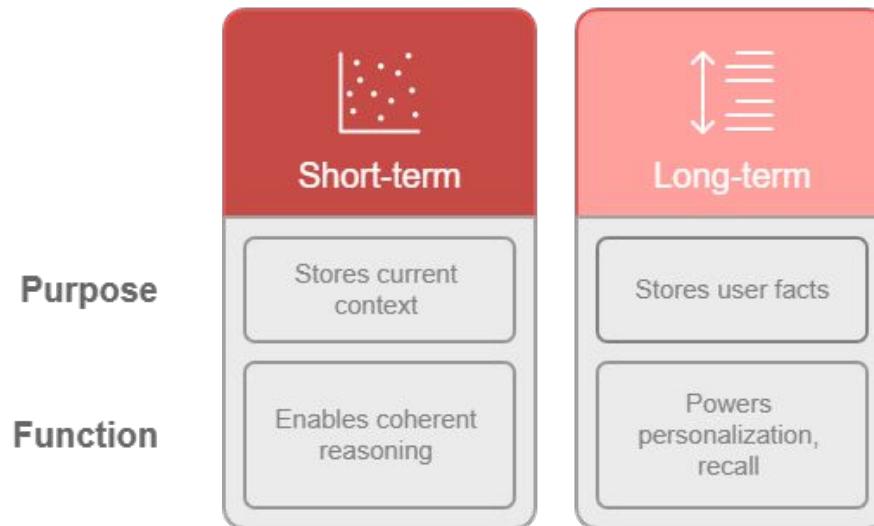
Component 5: Memory

- Stores **key info**:
 - internal thoughts, tool outputs, user data
- Retains **context** across steps and sessions
- Helps agent learn from **past interactions**
- Crucial for **personalization** and **long-term tasks**





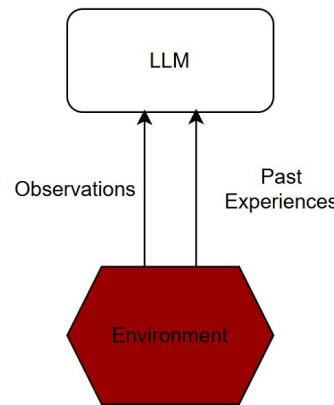
Types of Memory





Observations & Feedback

- Capture **what worked** and **what failed**
- Store **tool results**, **task completions**, or **user feedback**
- Enable agents to **refine** their strategy or escalate





Prompt Engineering as Policy

- Prompts can act like **behavioral blueprints** for the agent
- **System prompts = policies**
- Guide:
 - Role and tone
 - Constraints and permissions
 - Decision-making strategies
- **Structured prompting:**
 - XML, JSON modes, tool selection guidelines

You have access to the following tools:

- calculator: **for** math operations
- search: **for** current **or** factual info
- code_executor: **for** running Python code

When choosing a tool, respond **using this** JSON format:

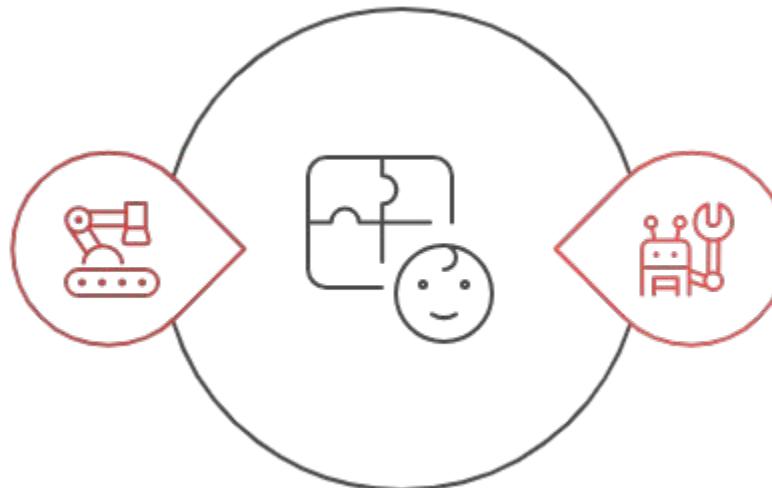
```
{  
    "tool": "tool_name",  
    "input": "input string"  
}
```



Orchestration Patterns in Agents

Multi-Agent System

Multiple specialized agents coordinating.
These agents collaborate, passing messages and refining each other's outputs.



Single-Agent System

A single model looping with tools and instructions



Multi-Agent Patterns



Manager Pattern

One agent coordinates and assigns tasks to others, promoting control and clarity.

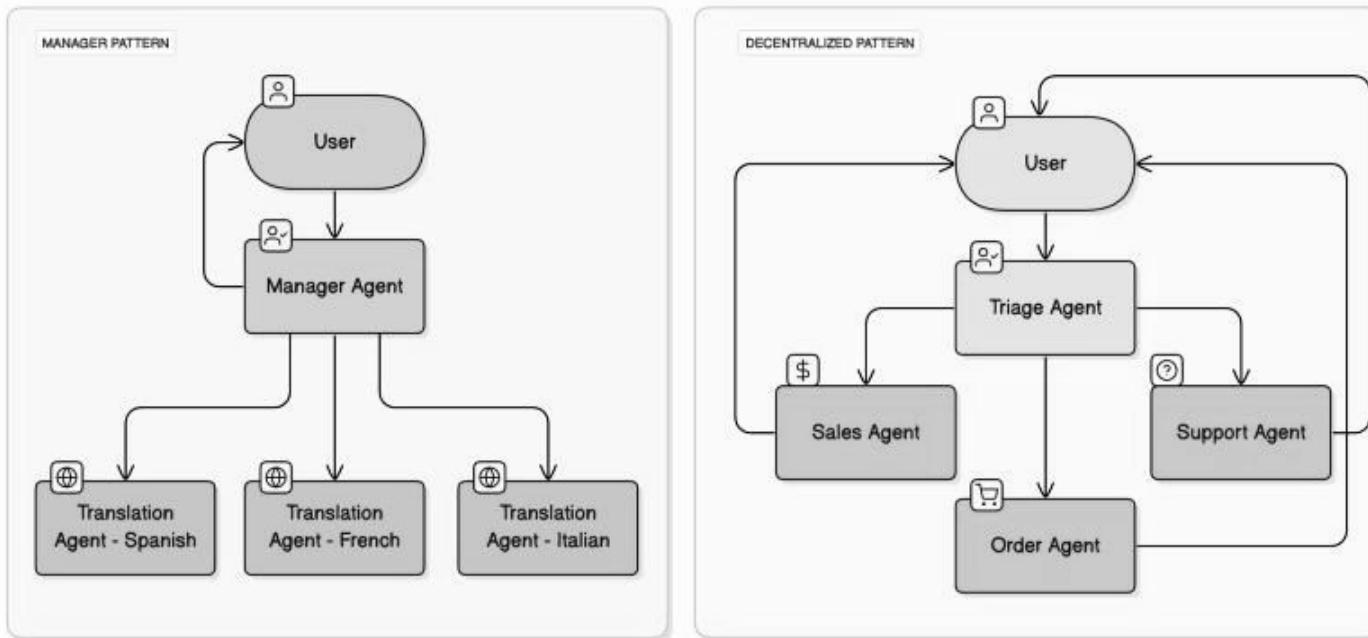


Decentralized Pattern

Agents communicate peer-to-peer without a central coordinator, enabling dynamic collaboration.



Multi-Agent Patterns





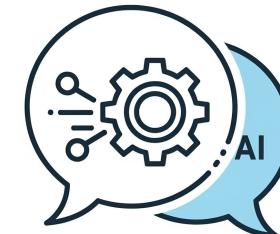
AMERICAN
UNIVERSITY
OF BEIRUT

Challenges of Agentic Systems



Poor Context Understanding

- Agents may **misinterpret** vague or nuanced instructions
- **Lack of grounding** in user goals or domain-specific language
- **Solution:**
 - Use **clarification prompts** before acting (e.g., “Can you confirm if you meant X or Y?”)
 - **Fine-tune** on domain-specific tasks or instruction-following datasets

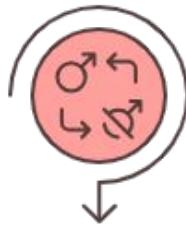




Hallucination and Reasoning Limits



LLMs may fabricate tools, steps, or results reason incorrectly or overcomplicate



Limited memory = loss of critical context



Poor self-evaluation



Tool Misuse and Failure Recovery

- Agents may call **incorrect** tools or **misuse** parameters
- Lack of **error handling** causes **failure** cascades
- **Example:** API returns 404 or malformed response → agent gets stuck
- **Solution:** Add tool schema validation, fallback options, retries, tool metadata





Observability and Debugging

- Debugging agents is **harder** than traditional scripts
- Agents involve:



Hidden
Reasoning



Long Prompt
Chains



Opaque
Decision-
Making

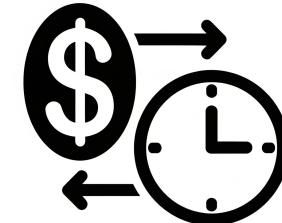


Multi-agent
Loops



Cost and Latency Trade-offs

- Tool use = more LLM calls = **more tokens**
- Planning + reflection = **longer latency**
- Agents may **retry, repeat, or loop** → unexpected cost spikes
- **Solution:** Use token limits, prompt summarization, and smart tool gating





AMERICAN
UNIVERSITY
OF BEIRUT

Conclusion



Future Directions



Enhanced Collaboration

Agents will work seamlessly with humans and other AI systems.



Contextual Intelligence

Improved understanding of nuanced and ambiguous instructions.



Autonomous Learning

Agents will adapt and optimize behavior with less human input.



Ethical and Transparent AI

Focus on fairness, explainability, and responsible decision-making.



Best Practices



Start with narrow tasks
and expand scope
incrementally



Use tool gating, max steps,
and cost monitoring



Log all steps and tool
calls for debugging and
transparency



Best Practices



Design clear prompt policies for consistent agent behavior



Include fallbacks and recovery strategies for safety



Continuously evaluate and fine-tune based on real user feedback



Agents are not just smarter chatbots — they are the foundation of a new way to automate thinking and doing



Hands_on: Exploring_and_Comparing_Large_Language_Model_Thinking_Strategies



AMERICAN
UNIVERSITY
OF BEIRUT

Thank you!