



AMERICAN  
UNIVERSITY  
OF BEIRUT

Fall 2025

# EECE 503P/798S: Agentic Systems

Chapter 8: Multi-Agent Systems




# Learning Objectives

- Explain when **multi-agent systems (MAS)** are preferable to **single-agent designs**
- Compare major **MAS architectures** and select an appropriate **topology** for a given task
- Specify **agent-to-agent (A2A)** messaging schemas and coordination protocols
- Design **memory strategies** for MAS (private vs shared vs mediated).
- Implement a minimal MAS with an **A2A protocol** and justify safety/guardrail choices



# Course Outline






AMERICAN  
UNIVERSITY  
OF BEIRUT

# When do we Need Multiple Agents



# Why do we need multiple agents?


- As the complexity of your tasks increases, **one ‘helpful assistant’ agent won’t be sufficient** to accurately execute your task.
- If your agent’s task is to create a fully functional multi-page website, you would need a planner, designer, programmer, quality assurance, and devops agents.
- The **quality of output would plummet** if you gave all the design, programming, and deployment tools to one agent





# Why do we need multiple agents?

## Agentic Design Patterns: Multi-Agent Collaboration




*Proposed ChatDev architecture. Image adapted from "Communicative Agents for Software Development," Qian et al. (2023).*



# Why do we need multiple agents?


- It might seem counterintuitive that, although we are making multiple calls to the same LLM, we apply the programming abstraction of using multiple agents.
- Even with long context windows, LLMs struggle with **complex inputs**.
- Agentic workflows improve performance by guiding the model to handle **one subtask at a time**, such as focusing on what matters when acting as a software engineer.





# Why do we need multiple agents?

- Multi-agent design breaks **complex tasks** into manageable **subtasks**.
- Similar to splitting a program into **processes or threads**.
- Makes building **complex systems** (e.g., a web browser) **easier**.





AMERICAN  
UNIVERSITY  
OF BEIRUT

# When to use and not to use Multi-Agent Systems



AMERICAN  
UNIVERSITY  
OF BEIRUT


# When should we use Multi-Agent Systems?



# When should we use Multi-Agent Systems?

## 1. Heterogeneous toolsets/permissions:

- Deciding on how to query a public dataset is very different from configuring an online search query.
- Having a **specialized agent for each tool** would allow you more room to carefully instruct the agents on how to use those tools.






# When should we use Multi-Agent Systems?

## 2. Tasks with natural pipelines

(ETL → analyze → generate  
→ verify → publish)






# When should we use Multi-Agent Systems?

## 3. Need for:


- **deliberation**  
(debate/critique),
- **redundancy** (ensembles),
- or **safety separation**  
(untrusted vs trusted zones).





# When should we use Multi-Agent Systems?

## 4. Multi-modal inputs or multi-domain expertise






AMERICAN  
UNIVERSITY  
OF BEIRUT

# When not to use Multi-Agent Systems?



# When should we not use Multi-Agent Systems?

- **orchestration complexity** outweighs **benefits**
- **latency/cost budgets** are tight





AMERICAN  
UNIVERSITY  
OF BEIRUT

# How do we Design a Multi-Agent System



# Pillars of a Multi-agent System

## *Multi-agent system*

### Division of Labor

→ Instead of having one agent with multiple distinct responsibilities, create a set of specialized set of agents per role.  
→ Example: Website developer agent, Website designer agent

### Checks and balances

→ Each of these agents have a distinct output. There should be checks in place to make sure that the agents' outputs are in the correct format and have executed without errors.  
→ All those agents append their knowledge and outputs into the memory. Thus, we should carefully manage appending and querying the memory.

### Parallelism

→ Not all tasks require sequential execution. For example, you can run multiple search agents in parallel.

### Reactivity


→ Agents perceive their environment and respond to changes in a timely manner.

### Proactiveness

→ Agents do not merely act in response to their environment; they are capable of taking initiative to fulfill their designed objectives.




# Multi-agent System Architectures





# Network Architecture - Swarm / Peer-to-Peer Mesh


- **Structure:** No central controller; agents share signals or stigmergic cues.
- **Pros:** Robust to failures; scalable via local rules; emergent behavior.
- **Cons:** Hard to debug; non-deterministic; needs careful termination design.
- **Use-cases:** Exploratory search; content curation; distributed monitoring.





# Network Architecture - Blackboard Architecture


- **Structure:** Agents read/write a **blackboard** (shared world state) and react to changes.
- **Pros:** Loose coupling; emergent cooperation; good for partial info fusion.
- **Cons:** Race conditions; requires arbitration/locks; provenance tracking needed.
- **Use-cases:** Multi-modal synthesis; incident response; analytics with many sources.





# Supervisor Architecture


- **Structure:** Central **Conductor** routes tasks among **Specialist** agents.
- **Pros:** Simple mental model; global visibility for guardrails; easy logging.
- **Cons:** Single bottleneck; conductor becomes complex; central point of failure.
- **Use-cases:** Customer-support triage; research-then-write flows; retrieval → drafting → QA.





# Supervisor Architecture - Agents as tools


- **Structure:** A single **Supervisor** plans and invokes **sub-agents** via *tool interfaces*. Each sub-agent exposes a narrow contract (input/output schema) and runs its own internal loop; the supervisor only observes structured I/O, not internal prompts.
- **Pros:** Strong isolation & testability; simple orchestration (tool-call API); per-tool auth & budgets; easy fallback/retries; compatible with mixed frameworks/models; unit-test sub-agents as black-box tools.
- **Cons:** Supervisor bottleneck; reduced shared context; risk of long tool chains; over-serialization (no peer-to-peer);





# Hierarchical Architecture - Planner


- **Structure:** Multi-level planning with **task decomposition** and escalation.
- **Pros:** Scales to large tasks; clear ownership; supports recursive refinement.
- **Cons:** Overhead in planning; risk of over-decomposition; requires termination rules.
- **Use-cases:** Long-horizon projects; enterprise workflows with approvals.





# Hierarchical Architecture - Marketer


- **Structure:** Manager announces task; Bidders propose; winner executes.
- **Pros:** Dynamic allocation; handles heterogeneous costs/skills.
- **Cons:** Auction overhead; strategy gaming; pricing alignment is nontrivial.
- **Use-cases:** Tool/time-cost-aware routing; resource-bounded compute.





# Hierarchical Architecture - Agents as tools

- **Structure:** Fixed sequence/graph of roles; outputs become next inputs.
- **Pros:** Deterministic, testable; parallel stages possible; great for SLAs.
- **Cons:** Rigid; backtracking requires side-channels; error propagation.
- **Use-cases:** Data ingestion & enrichment; report generation; code CI assistants.






AMERICAN  
UNIVERSITY  
OF BEIRUT

# Communication between Multiple Agents



# Agent 2 Agent Protocol

- Launched by Google in April 2025, the Agent-to-Agent protocol is designed to enable **seamless interoperability within multi-agent systems**.
- It allows a user-facing "client" agent to **orchestrate** and delegate tasks to **specialized "remote" agents**.
- The protocol is built on standard web technologies like HTTP, JSON-RPC, and Server-Sent Events (SSE).






# Agent 2 Agent Protocol - Components

1- **Agent Card:** Each remote agent publishes a standardized JSON file, or "Agent Card," that advertises its capabilities, skills, and authentication requirements. This allows client agents to discover the right specialist for a task.

## Agent Card



Name: Technical Writer


Description: Writes comprehensive blogs on technical topics

Skills: Breaks down complex topics into accessible writing



# Agent 2 Agent Protocol - Components


**2- Task:** The central unit of work. When a client agent sends a request, it creates a task that moves through a well-defined lifecycle (e.g., submitted, working, completed, input-required).





# Agent 2 Agent Protocol - Components


**3- Messaging:** Agents communicate by exchanging messages. A single message can contain one or more "parts" for different content types, such as text, structured data (JSON), or files.






# Agent 2 Agent Protocol - Components

4- **Artifacts**: Immutable outputs or results generated by a remote agent upon completing a task.



# Agent 2 Agent Protocol



1

2

3

4

5

## Receive Request

Client agent receives user request

## Select Agent

Client agent chooses suitable remote agent

## Send Task

Client agent sends task to remote agent

## Process Task

Remote agent processes the task


## Return Result

Remote agent sends back the completed task



# MCP vs A2A

## MCP vs. A2A protocol





AMERICAN  
UNIVERSITY  
OF BEIRUT

# Hands on (Architectures + A2A): `1_multi_agent_architectures.ipynb`




AMERICAN  
UNIVERSITY  
OF BEIRUT

# How do Agents Share Memory




# Multi-agent shared memory




# Collaborative Memory Framework (Rezazadeh et al., 2025)

- Introduces **two-tier memory**: private memory (user-specific) and shared memory (selectively shared).
- Uses **bipartite graphs** to encode asymmetric, time-evolving access permissions.
- Immutable provenance metadata** ensures **auditability** and **retrospective permission checks**.



# Memory Sharing in Practice: Real-Time and Persistent Memory


- Memory-Sharing (MS) framework aggregates **query-response pairs** from multiple agents.
- Enables agents to **retrieve relevant examples** from a **shared memory pool** to improve open-ended task performance.
- Empirically shown to **enhance reasoning** and **knowledge transfer** in specialized domains.





# Anthropic's Research Agent Memory

## High-level Architecture of Advanced Research






AMERICAN  
UNIVERSITY  
OF BEIRUT

# Challenges of Multi-Agent Systems




# Challenges of multi-agent systems


- Instruction **leakage** between roles; prompt **drift**; duplicated work; **infinite loops**.
- **Inconsistent schemas**; brittle coupling to tool **I/O**; **silent failures** on the bus.
- **Grounding gaps**: agents cite each other rather than sources.



a. Equi-level structure



b. Hierarchical structure



c. Nested structure



# Conclusion

- Agents are stateful and errors compound.
- **Agents can run for long periods of time, maintaining state across many tool calls.** This means we need to durably execute code and handle errors along the way.
- Without effective mitigations, **minor system failures can be catastrophic for agents.**
- When errors occur, we can't just restart from the beginning: **restarts are expensive** and frustrating for users.
- Instead, we built systems that can **resume from where the agent was when the errors occurred.**



AMERICAN  
UNIVERSITY  
OF BEIRUT

# Thank you